

Lecture 3: Introduction to Point-based Graphics

Hamid Laga

Global Edge Institute
Tokyo Institute of Technology
hamid@img.cs.titech.ac.jp
<http://www.img.cs.titech.ac.jp/~hamid/>



Remark

- This lecture is a summary of Siggraph 2004 course on Point-based Computer Graphics:
 - <http://graphics.stanford.edu/~mapauly/Pdfs/SigCourse04.pdf>

2



Overview

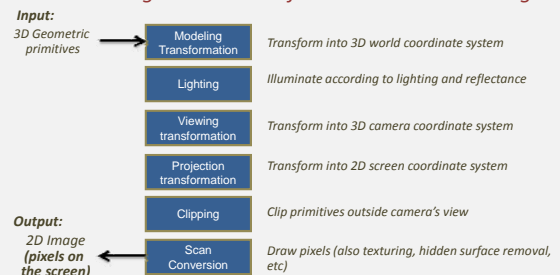
- Introduction and motivation
 - The rendering pipeline
 - Representations in graphics revisited
 - Volume, triangle → points
- 3D geometry acquisition
 - 3D scanners technology
- Point-based representation
 - Surface elements: *Surfel*
- Processing and editing
 - (skip)
- Rendering
- Summary

3




The rendering pipeline

3D Rendering: Conversion of a 3D scene into an image

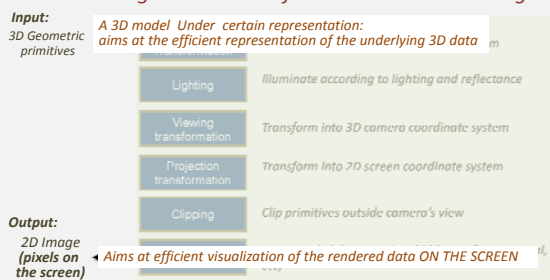


4



The rendering pipeline

3D Rendering: Conversion of a 3D scene into an image




5



Representation in Graphics

- Volume representation
 - Extension of 2D images into 3D images
 - 3D grid of cells (voxels)
 - A voxel stores some properties of the model
 - Suitable when we want to represent interior details:
 - Medical applications,
 - Fluid simulation,
 - Natural phenomena simulation,
 - Etc...




6



Representation in Graphics

- **Volume representation**
 - 3D grid of cells called *voxels*
 - Each voxel stores some properties of the model
 - Extension of 2D images into 3D images
 - Suitable when we want to represent interior details:
 - Medical applications,
 - Fluid simulation,
 - Natural phenomena simulation,
 - Etc,...



7

Representation in Graphics

- **Surface representation**
 - Represents the boundaries of the 3D object:
 - Parametric (ex. Splines, NURBS)
 - implicit representations
 - Polygon-based representations
 - and Point-based representations

8


Representation in Graphics

- **Parametric and implicit representations**
 - **Properties**
 - ✓ Rigorous mathematical concept
 - ✓ Robust evaluation of geometric entities
 - ✓ Shape control for smooth shapes
 - ✗ Require proper parameterization
 - ✗ Discontinuity modeling
 - ✗ Topological flexibility

9

Polygon-based representations

- **Functions → Triangles**
 - Piecewise linear approximations:
 - Given a continuous surface S , it is approximated by a set of triangles



10

Polygon-based representations

- **Functions → Triangles**
 - Piecewise linear approximations:
 - Given a continuous surface S , it is approximated by a set of triangles
 - Irregular sampling of the surface
 - No parameterization needed (for now)

11


Polygon-based representations

- **Polynomial → Triangles**
 - **Properties**
 - ✓ Simple and efficient representation
 - ✓ Hardware pipelines support Δ
 - ✓ Advanced geometric processing
 - ✓ The widely accepted queen of graphics primitives
 - ✗ Sophisticated modeling is difficult
 - ✗ (Local) parameterizations still needed
 - ✗ Complex LOD management
 - ✗ Compression and streaming is highly non-trivial

12

Point-based representation

- **Triangles → Points: remove connectivity!!**
 - Piecewise linear functions to Delta distributions
 - Discrete samples of geometry
 - No connectivity or topology – most simple
 - Store all attributes per surface sample



13

GLOBAL EDGE INSTITUTE

Point-based representation

- **Triangles → Points: remove connectivity!!**
 - Piecewise linear functions to Delta distributions
 - Discrete samples of geometry
 - No connectivity or topology – most simple
 - Store all attributes per surface sample
 - Point clouds
 - Points are natural representations for 3D acquisition systems
 - Meshes constitute an enhancement of point samples

14

GLOBAL EDGE INSTITUTE

Motivations (1/2)

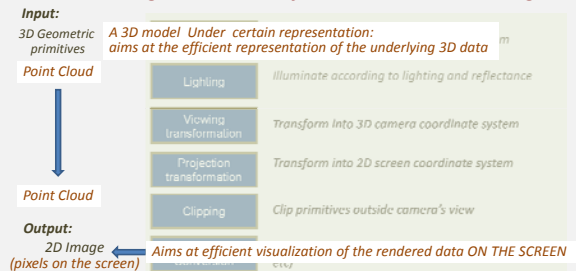
- **Polygonal complexity**
 - Complex models require very large set of polygons
 - Accurate representation requires more than 200,000 polygons
 - Carrying out the connectivity is useless because:
 - Size of polygon \ll size of a pixel
 - Storage requirement is very expensive

15

GLOBAL EDGE INSTITUTE

The rendering pipeline revisited

3D Rendering: Conversion of a 3D scene into an image



16

GLOBAL EDGE INSTITUTE


Motivations (2/2)

- **Upcoming 3D digital photography**
 - 3D scanners provide point clouds
 - Generating a consistent triangle mesh is time consuming and difficult
 - Points represent organic models (feathers, trees) much more readily than polygons

17

GLOBAL EDGE INSTITUTE

3D Content Creation Pipeline



Points as a generalization of pixels !!

18

GLOBAL EDGE INSTITUTE

Overview

- Introduction and motivation
 - The rendering pipeline
 - Representations in graphics revisited
 - Volume, triangle → points
- 3D geometry acquisition
 - 3D scanners technology
- Point-based representation
 - Surface elements: *Surfel*
- Processing and editing
 - (skip)
- Rendering
- Summary

19

Acquisition of geometry and appearance

How can we capture reality ?

- Real objects have complex geometry
 - feather, bonsai trees, ...
- Appearance in arbitrary environments with arbitrary lighting
 - rendering a glass of beer in a complex surroundings
 - Specular, fuzzy, and transparent materials are very difficult to scan.



20

Acquisition of geometry and appearance

- Goals
 - Fully-automated 3D model creations
 - (or minimum user input)
 - Faithful acquisition and representation of appearance
 - Good geometry is always needed
 - Faithful reproduction of the appearance.
 - Placement into new virtual environments



21

3D geometry acquisition

- Contact digitizers
 - Direct physical contact with the object to scan
 - Lot of manual labor
- Passive methods (using cameras)
 - Requires that the object has texture
 - Requires that the BRDF of the objects is Lambertian
- Active light imaging systems
 - Use active laser range scanners
 - Very accurate
 - Severe restrictions on the types of materials
 - Ex.: specular and fuzzy materials are difficult.



4 millions points
[Levoy et al. 2000]


22

Stereo-vision for 3D acquisition

23

Binocular vision

- Why humans have two eyes ?
 - Because of our binocular vision we can
 - See **WHERE** objects are in relation to our own bodies with much greater precision,
 - See a little bit **AROUND** solid objects without moving our heads,
 - Perceive and measure **EMPTY** space around us.



Source: <http://www.vision3d.com>


24

Binocular Vision Blessings

- If You've Got Stereo Vision, Count Your Blessings!
 - The American Academy of Ophthalmology, Sept., 1996:
 - "many occupations are not open to people who have good vision in one eye only [i.e. people without stereo vision]"
 - Few examples of occupations that depend heavily on stereo vision:
 - Baseball player
 - Waitress, Driver
 - Architect, Surgeon, Dentist
 - Few examples of general actions that depend heavily on stereo vision:
 - Throwing, catching or hitting a ball
 - Driving and parking a car
 - Planning and building a three-dimensional object
 - Threading a needle and sewing
 - Reaching out to shake someone's hand
 - Pouring into a container


Stereo Vision

- Example of stereo cameras





Finding the 3D position

- 3D reconstruction (*inverse imaging process*)
 - Given a point in the 2D image, what are its 3D coordinates in the 3D space





Finding the 3D position

- 3D reconstruction (*inverse imaging process*)
 - Given only a point in the 2D image, what are its 3D coordinates in the 3D space



Finding the 3D position

- 3D reconstruction (*inverse imaging process*)
 - Given only a point in the 2D image, what are its 3D coordinates in the 3D space




Finding the 3D position

- 3D reconstruction (*inverse imaging process*)
 - Given a point in the 2D image, what are its 3D coordinates in the 3D space

$$\begin{bmatrix} X \\ Y \\ Z \\ 1 \end{bmatrix} = \frac{1}{C} \begin{bmatrix} * & * & * & * \\ * & * & * & * \\ * & * & * & * \\ * & * & * & * \end{bmatrix}^{-1} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix}$$

P⁻¹: inverse of 3 X 4 matrix !!!!!




The projection ambiguity

Finding the 3D position

- Incorporate other cues to resolve the projection ambiguity:
 - A priori knowledge, high-level semantic knowledge
- For humans
 - Monocular cues
 - Relative size between objects
 - Familiar sizes
 - Motion parallax, vanishing points, etc...
 - Eye movements
 - Binocular cues
 - Stereo-vision: this is the most developed for normal people
 - 3D Shape from silhouettes

31


Example of stereo camera



32

Stereo vision: how does it work

- A flavor of Projective Geometry



33

Stereo vision: how does it work





Figure 9.8 Epipolar geometry in stereopsis.

34

Stereo vision: towards 3D perception

- Simplified case:
 - C1 and C2 are horizontal (the case of BumbleBee 2)



35

Geometry of a stereo camera

- Three main targets:
 - Scene geometry (3D reconstruction)
 - Given:
 - a set of corresponding points $\{x_i \leftrightarrow x'_i\}$,
 - the cameras P and P'?
 - Can we find the 3D point X?
 - The answer is YES
 - The stereo-camera allows to compute the depth of each point in the image
 - [DEMO WITH BUMBLEBEE@2]

36

Geometry of a stereo camera

- **Properties**
 - The state of the art of stereo methods generate noisy depth maps
 - It works better on low resolution images (320x240)
 - Many other parameters to tune
- [Illustrate these on the BumbleBee demo]

37

Multiple view geometry


- **Why not using more than two cameras!!**



38

Multiple view geometry

- **Networking multiple camera systems**



39

Multiple view geometry

- **Networking multiple camera systems**
 - Cover a wide area
 - Deal with occlusions
 - Accurate 3D reconstruction
 - Application to video surveillance

40

3D Shape from silhouette

- Stereo vision is not the only way for building the 3D structure of the world.
- [SHOW A VIDEO]


41

Appearance acquisition

- **BRDF estimation, inverse rendering**
 - Fit parametric BRDF model to the observed measurements
 - BRDF models are not capable to represent complex materials (feathers, hair, clothes, etc)
 - Cannot capture inter-reflections, self-shadowing and subsurface scattering
- **Image-based modeling + Point-based rendering**
 - Images represent objects regardless the complexity of its geometry and appearance
 - Cost to pay:
 - Need of more data

42

Example of acquisition system



[matsushik et al. 2002] *Image-based 3D photography using opacity hulls*, in SIGGRAPH 2002


43

Example of acquisition system

- **General idea**
 - Take multiple pictures of the object from multiple views
- **System architecture**
 - Objects place onto a plasma display, mounted on a rotating platform
 - Array of lights mounted on a rotating overhead platform
 - Six video cameras on a stationary arc, pointed toward the object
 - Plasma displays:
 - For multi-background matting techniques to acquire alpha mattes of the object from multiple viewpoints
 - Estimate light transport through transparent and reflective objects
 - They are switched off after the alpha matte acquisition

44

System overview




45

Example of acquisition system

- **Output of the system**
 - Point cloud approximating the geometry of the shape
 - Opacity lightfield at each surface point
 - Think of RGBA values
 - It captures reflections and transparency
 - In general case:
 - Each surface point carries some properties, simplest case is when the point is carrying just its color.

46

Example of results



47

Example of acquisition system

- **Many steps are required in this system**
 - Acquisition (we already talked a little bit)
 - Model
 - **Point-based rendering**
 - Relighting
 - Reflectance field estimation

48

Overview

- Introduction and motivation
 - The rendering pipeline
 - Representations in graphics revisited
 - Volume, triangle → points
- 3D geometry acquisition
 - 3D scanners technology
- Point-based representation
 - Surface elements: *Surfel*
- Processing and editing
 - (skip)
- Rendering
- Summary

49

Points as rendering primitives

- Point clouds instead of triangle meshes
 - Introduced in 1985 by [Levoy and Whited 1985]
 - Analogy in 2D:
 - 2D vector graphics vs. pixel graphics



Triangle mesh with texture




Point cloud

50

Surface elements - Surfels



- Each point corresponds to a surface element, or **surfel**, describing the surface in small neighborhood
- Basic surfels:
 - A point with position and a constant color



51

Surface elements - Surfels

- In general
 - Surfels can be extended to store additional attributes
 - The union of surfel discs covers the object such that there are no holes on the surface




52

Overview

- Introduction and motivation
 - The rendering pipeline
 - Representations in graphics revisited
 - Volume, triangle → points
- 3D geometry acquisition
 - 3D scanners technology
- Point-based representation
 - Surface elements: *Surfel*
- Processing and editing
 - (skip)
- Rendering
- Summary

53


Point rendering pipeline



- Simple, pure forward mapping pipeline
- Surfels carry all information through the pipeline

54


Point rendering pipeline



- **Shading:**
 - Input surface samples are shaded according to the measured RGBA values
- **Projection (warping)**
 - Points are projected onto the image space.
- **Continuous surface reconstruction**
 - A continuous surface is reconstructed in the image space (similar to triangle rasterization)
- **Filtering / sampling**
 - Filtering and sampling the continuous representation at the output pixel positions (similar to texture filtering and sampling)

55


Point rendering pipeline




- **Shading**
 - The scanner acquires points and colors from different (pre-set) viewing points
 - We want to assign a color to point when rendered from a NEW view point:
 - Blend colors based on the angle between new viewpoint and acquired RGBA images
 - Linear combination of colors from closest views

56

Point rendering pipeline




- **Continuous reconstruction**
 - Similar to triangle rasterization.
 - We want to do similar thing for surfels when projected to the image domain.



57

Reconstruction



- **Goal**
 - Estimate surface properties at unknown locations



58

Reconstruction


- **Step 1**
 - Define local coordinate frame at each point
 - defined by the normal vector
 - Define a reconstruction kernel associated with each point in the local coord. frame:
 - Usually a Gaussian Kernel



59

Reconstruction


- **Step 3**
 - Mapping into a global parameter domain
 - For each reference frame we define a parameter mapping that convert local coordinates to coordinates in the global parameter domain



60

Reconstruction

- Step 3
 - Mapping into a global parameter domain
 - For each reference frame we define a parameter mapping that converts local coordinates to coordinates in the global parameter domain



61

Reconstruction

- Step 4
 - Evaluate the reconstruction equation to reconstruct the surface

$$g(x) = \frac{\sum_i f(x_i) r_i(\phi_i^{-1}(x))}{\sum_i r_i(\phi_i^{-1}(x))}$$


Parameterization, i.e., mapping to global parameter domain

- Phi is the mapping from local coord frame to global coord frame.

62

Reconstruction


- Step 4
 - Evaluate the reconstruction equation to reconstruct the surface



63


Reconstruction

- Step 4
 - In 3D



64

Point rendering pipeline



- Projection / warp
 - Exactly the same as polygon representation

65


Warping

- Steps:
 - Use image space as global parameter domain
 - Warp local reference domains to image space using 2D to 2D projective mapping
 - Compute weighted sum of reconstruction kernels in image space

66

Warping

The pose of the tangent plane and its perspective projection determine the 2D-2D mapping



67

Warping

- Reconstruction kernels in image space are called *footprints* or *splats*

$$g(j, k) = \frac{\sum_i f_i r_i(\phi_i^{-1}(j, k))}{\sum_i r_i(\phi_i^{-1}(j, k))}$$

output pixel coordinates perspective projection of reference domain i

Illustration of the reconstruction of the color of a point

68

Splatting algorithm


- The rendering algorithm is called *splatting algorithm*

```
for each sample point {
  shade surface sample;
  splat = projected reconstruction
  kernel;
  rasterize and accumulate splat;
}
for each output pixel {
  normalize;
}
```

69


Splatting algorithm

- Visibility: *Hidden surface removal*



70

Point rendering pipeline




- Filtering / resampling
 - Sample the reconstructed signal at the output pixel locations:
 - Apply some filtering to the reconstructed signal before sampling it (to achieve high image quality)

71

Filtering / resampling

- Antialiasing
 - Rendering using warped reconstruction kernels does not prevent aliasing artifacts:



Moire patterns on highly structured patterns



Disintegration of textures with highly fine details

72

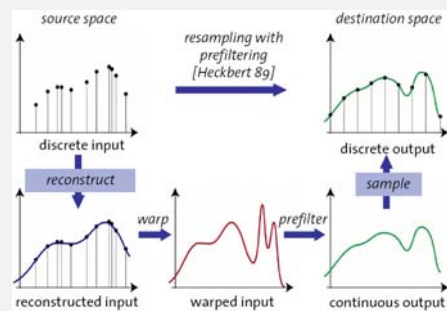
Filtering / resampling

- **Antialiasing**

- Rendering using warped reconstruction kernels does not prevent aliasing artifacts:
 - Aliasing occurs when sampling a signal that is not band-limited to the Nyquist frequency of the sampling grid
 - Aliasing is avoided by pre-filtering, i.e., band-limiting to the Nyquist frequency before sampling
- To avoid aliasing artifacts:
 - Re-sampling with pre-filtering

73

Resampling with pre-filtering



74

Rendering results

Head 429,075 points Helicopter 987,552 points Matterhorn 4,782,011 points



75

Performance

- Software implementation [Zwicker et al. 2001]
 - 1.1 GHz AMD Athlon, 1.5 Gbyte RAM

Data	#Points	256 x 256	512 x 512
Scanned head	429,075	1.3 fps	0.7 fps
Helicopter	987,552	0.6 fps	0.3 fps
Matterhorn	4,782,011	0.2 fps	0.1 fps



Scanned head Helicopter Matterhorn

76

Applications of point rendering

- Direct visualization of point clouds from 3D scanners [Rusinkiewicz et al. 2000, Matusik et al. 2002, Xu et al. 2004]
- Real-time 3D reconstruction and rendering for virtual reality applications [Gross et al. 2003]



[Rusinkiewicz et al. 2000]



[Matusik et al. 2002]



[Gross et al. 2003]

77

Applications of point rendering

- Hybrid point and polygon rendering systems [Chen et al. 2001, Cohen et al. 2001]
- Rendering animated scenes [Wand et al. 2002]



[Chen et al. 2001]



[Wand et al. 2002]


78

Applications of point rendering

- Interactive display of huge meshes [Wand et al. 2001, Deussen et al. 2002]
- On the fly sampling and rendering of procedural objects [Stamminger et al. 2001]
- Modeling and editing [Zwicker et al. 2002, Pauly et al. 2003]



[Deussen et al. 2002]



[Stamminger et al. 2001]



[Pauly et al. 2003]

79

Conclusions Summary

- **Point rendering as a re-sampling process**
 - High quality point splatting with antialiasing using a pre-filtering approach
- **Performance**
 - 3 million high quality splats per second with hardware support [Zwicker et al. 2004]
 - 1 million points per second in software
- **Various applications for point splatting have been explored**

80

Future work

Interesting directions to explore

- Rendering hardware
 - Dedicated point rendering hardware
- Extension of current hardware
- **Unifying rendering primitives for all sorts of surface representations**
 - Rendering architecture for on the fly sampling and rendering

81

Resources

- **Online resources:**
 - Point-based graphics resource:
<http://hksmyweb.hinet.net/Resource/PointBasedPaper.html>
- **Online courses:**
 - Siggraph 2004 course:
<http://graphics.stanford.edu/~mapauly/publications>
- **Software**
 - PointShop3D: (Computer Graphics Lab at ETH Zurich)
<http://www.pointshop3d.com>
 - Qsplat (Computer Graphics Lab at Stanford)
<http://graphics.stanford.edu/software/qsplat/>

82

Don't forget

- **Lecture's slides:**
 - Check my homepage to download the latest version of the lecture slides:
 - <http://www.img.cs.titech.ac.jp/~hamid/courses/>

83

Questions

84